

Visesh Bentula

(737) 354-0646 | visesh66@gmail.com

linkedin.com/in/viseshb | github.com/viseshb | viseshb.github.io/My-Portfolio

PROFESSIONAL SUMMARY

Software Engineer with nearly two years of industry experience building AI systems, backend platforms, and LLM infrastructure. Experienced in agent orchestration, RAG pipelines, data security, and full-stack development. Launched VibeQuorum, an AI platform that attracted 300+ organic waitlist signups. Master of Science in Computer Science with a focus on AI and machine learning from Texas A&M University.

EDUCATION

Texas A&M University | Master of Science in Computer Science (GPA: 3.8/4.0), San Antonio, USA | Jan 2025 - May 2026

- Research Assistant, Deep Learning: Conducted research on reducing hallucinations in RAG systems by testing transformer attention-based methods and evaluating generated answers with RAGAS for faithfulness and answer relevancy.
- Coursework: Artificial Intelligence, Machine Learning and Deep Learning, Algorithms, Operating Systems, Secure Software Development.

EXPERIENCE

Senior Software Engineer | Digifyde, Remote, USA | Jun 2026 - Present

- Built the Agent Harness for Orion, an enterprise data platform, creating a centralized system to manage LLM agents, prompts, tools, and workflows across AI-powered features.
- Led data mapping for Orion and built PII redaction and filtering guardrails that prevent sensitive client data from reaching external LLMs, with access protected through role-based authentication and authorization.
- Built full-stack AI features for Orion by organizing client data through a medallion pipeline across PostgreSQL and MongoDB, while working directly with clients to refine prompts and improve the accuracy of AI-generated results.

Software Engineer | Thermal Systems, Hyderabad, India | Jan 2024 - Dec 2024

- Reduced ERP reporting latency across 16 departments from 8 seconds to 350 milliseconds over 850K+ daily records by analyzing PostgreSQL execution plans, adding composite indexes, rewriting expensive joins, partitioning high-volume tables, and caching repeated queries in Redis.
- Increased production ML inference API throughput by 45% by containerizing services with Docker on AWS, caching repeated requests in Redis, batching model inputs, and horizontally scaling instances based on traffic.
- Reduced Sev-1 incident recovery time for Kafka-based microservices from 4 hours to under 30 minutes by adding Prometheus and Grafana alerts, centralized ELK logging, Kubernetes health checks, automatic pod restarts, and failure recovery workflows.

Software Engineer, AI and Data Platforms | Syra Health, Hyderabad, India | Aug 2023 - Jan 2024

- Built real-time healthcare data pipelines using Kafka, PostgreSQL, Elasticsearch, and GraphQL to ingest, process, and serve patient data for risk-assessment and fraud-detection models supporting 220K+ users.
- Reduced release time by 40% through GitHub Actions and Docker CI/CD pipelines with zero-downtime deployments and automated rollbacks, while React dashboards for healthcare risk and fraud insights increased user engagement by 20%.

PROJECTS

VibeQuorum | Agentic SaaS | LangGraph | Cloud Run | vibequorum.app

- Launched VibeQuorum, a multi-agent AI platform that helps solo builders and vibe coders review product architecture, tech stacks, API designs, security risks, implementation plans, and estimated costs, attracting 300+ organic waitlist signups without any marketing.
- Architected and shipped VibeQuorum's multi-agent backend with FastAPI and LangGraph, supporting 12 specialized agents through 68 endpoints and integrating 15+ LLM providers, OAuth, Stripe, Redis, Supabase, and Cloud Run.
- Validated the multi-agent orchestration layer through 400+ logged debates and 2,400+ cross-provider LLM calls across Google, AWS Bedrock, and NVIDIA, maintaining a sub-5% workflow failure rate before scaling the production infrastructure.

Context-Aware AI Assistant | MCP Orchestration | Tool Routing | Open Source | github.com/viseshb/context-aware-ai-assistant

- Built an AI assistant using FastAPI and Next.js that lets engineers search GitHub, Slack, and PostgreSQL through natural-language queries, reducing incident investigation time from 30-60 minutes to under 30 seconds using 16 MCP tools, source attribution, sub-400ms routing, JWT authentication, RBAC authorization, rate limiting, and automatic PII redaction.

SKILLS

Languages: Python, JavaScript, TypeScript, SQL, Java, C++

Frontend: React, Next.js, TypeScript, Tailwind CSS, Electron, HTML, CSS

Backend and Systems: FastAPI, Node.js, Express.js, REST APIs, GraphQL, Microservices, System Design, WebSockets, Webhooks

Data and Databases: PostgreSQL, MongoDB, Redis, Elasticsearch, Kafka, pgvector, Medallion Architecture (Bronze, Silver, Gold)

Infrastructure and Cloud: Docker, Kubernetes, GCP (Cloud Run), AWS (EC2, EKS, Lambda, Bedrock), Supabase, CI/CD (GitHub Actions), Prometheus, Grafana, ELK

AI, Security, and Identity: PyTorch, LLMs, RAG Pipelines, RAGAS, Multi-Agent Systems, Agent Orchestration, LangGraph, Vector Databases, Model Context Protocol (MCP), Prompt Engineering, OpenAI, Anthropic, OpenRouter, PII Redaction, Data Guardrails, JWT, OAuth, RBAC