

Visesh Bentula

+1(737)-354-0646 | visesh66@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

SUMMARY

Software Engineer with 2 years of industry experience building agentic systems, LLM infrastructure, and production backends. Founded and solo-engineered VibeQuorum, a multi-agent SaaS with 300+ waitlist signups. MS in Computer Science (AI/ML) from Texas A&M University, on OPT starting June 2026.

EDUCATION

Texas A&M University | **Master's in Computer Science (GPA: 3.8/4.0)**, San Antonio, USA

Jan 2025 – May 2026

- **Research Assistant, Deep Learning:** Investigating hallucination mitigation in retrieval-augmented generation (RAG) pipelines using transformer attention mechanisms for factually consistent large language model outputs.
- **Coursework:** Artificial Intelligence, Machine Learning & Deep Learning, Algorithms, Operating Systems, Secure Software Development

EXPERIENCE

Software Engineer | **Thermal Systems**, Hyderabad, India

Jan 2024 – Dec 2024

- Optimized SQL query latency from 8s to 350ms by redesigning PostgreSQL schemas, indexes, and query execution plans, sustaining low-latency analytics and ML inference over 850K+ records/day.
- Improved ML model-serving API throughput 45% by containerizing microservices with Docker on AWS, implementing Redis caching and request batching to eliminate redundant computation, and enabling horizontal auto-scaling to sustain consistent inference performance under peak traffic loads.
- Reduced mean time to resolution (MTTR) from 4 hours to under 30 minutes for Sev-1 production incidents by building real-time alerting with Prometheus and Grafana, structured logging via the ELK stack, and Kubernetes-based self-healing orchestration with async processing using Kafka across distributed services.

Software Developer – AI & Data Platforms | **Syra Health**, Hyderabad, India

Aug 2023 – Jan 2024

- Built real-time healthcare data pipelines with a 4-person engineering team using Kafka for streaming ingestion, PostgreSQL for structured storage, and Elasticsearch for low-latency search, exposing data through GraphQL APIs to power predictive models for health risk assessment and insurance fraud detection across a platform serving 220K+ users.
- Accelerated release cycles by 40% by building CI/CD pipelines with GitHub Actions and Docker, enabling zero-downtime deployments, and boosted user engagement 20% by shipping React dashboards with Figma and Framer API integrations for real-time analytics visualization.

Software Development Intern | **RADcube**, Hyderabad, India

Jan 2023 – Apr 2023

- Built REST APIs using Java and Spring Boot with JUnit test coverage, developed MERN stack components for end-to-end data flow, and authored technical documentation to streamline onboarding and ramp-up time for new developers in an Agile workflow.

PROJECTS

VibeQuorum | **Founder & Sole Engineer** | **agentic SaaS** · **LangGraph** · **Cloud Run** | [vibequorum.app](#)

- Built a production SaaS on FastAPI, LangGraph, and Next.js 15 where a tiered agent council (4 to 10 agents across 4 subscription tiers) debates schema, API contracts, security posture, and infrastructure cost for 5 to 30 minutes to produce a complete technical spec covering stack decisions, database migrations, CI/CD pipelines, and deployment from a single natural-language brief.
- Engineered a per-user OpenRouter provisioning pipeline where each signup auto-allocates isolated API credits sized by subscription tier with monthly rollover, routes agent workloads exclusively to frontier models (GPT, Claude, Gemini), and deployed the 46.8K-line codebase (26.7K Python, 20.1K frontend/config) on Google Cloud Run with Supabase (database) and Upstash Redis (cache), reaching 300+ waitlist signups at launch.

Context-Aware AI Assistant | **Open Source** | **MCP orchestration** · **tool routing** | [GitHub](#)

- Built an AI assistant on FastAPI and Next.js that lets engineers query GitHub, Slack, and PostgreSQL in plain English, cutting cross-tool incident searches from 30–60 minutes down to under 15 seconds by routing 16 MCP tools (7 GitHub, 5 Slack, 4 PostgreSQL) across 3 data sources with full source attribution on every answer.
- Implemented 8-layer security (JWT, RBAC, rate limiting, PII auto-redaction, read-only SQL enforcement, allowlists, input validation, audit logging) with sub-400ms tool dispatch, 5ms PII scan overhead, and 10 switchable LLM providers across 100+ concurrent WebSocket connections.

SKILLS

Core: Python, Java, JavaScript, TypeScript, SQL, Git, Linux, System Design, Distributed Systems

Systems / Backend: FastAPI, Spring Boot, Node.js, Express.js, Next.js, REST APIs, GraphQL, Microservices, PostgreSQL, MySQL, Redis, Kafka, Elasticsearch, WebSockets, Async Processing

Infrastructure & MLOps: Docker, Kubernetes, AWS (EC2, S3, RDS), Terraform, CI/CD (GitHub Actions), ELK Stack, Prometheus, Grafana, ONNX Runtime, Zero-Downtime Deployment

AI/ML: PyTorch, Hugging Face Transformers, LLMs, Agentic AI, Multi-Agent Systems, LLM Orchestration, LangGraph, RAG Pipelines, Embeddings, Vector Databases (pgvector), MCP Orchestration, Prompt Engineering, Model Evaluation

CERTIFICATIONS

[Google AI Professional Certificate](#) – Google (Coursera)

Apr 2026

[Machine Learning Specialization](#) – DeepLearning.AI & Stanford Online (Coursera)

Dec 2025

[Advanced SQL](#) – Great Learning Academy

Jun 2024